

Index SQL

Born to B-Tree

WHY !?



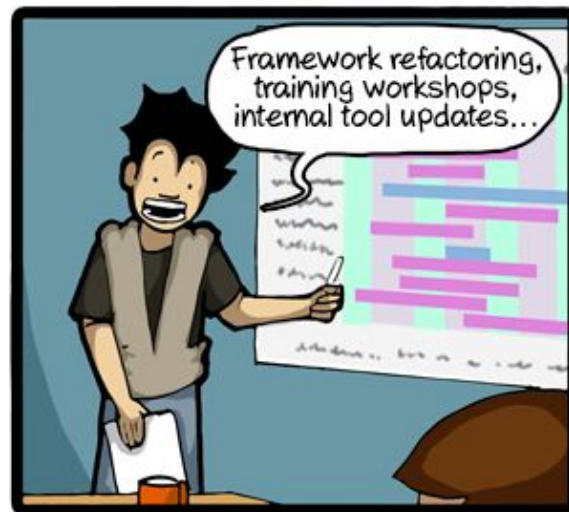
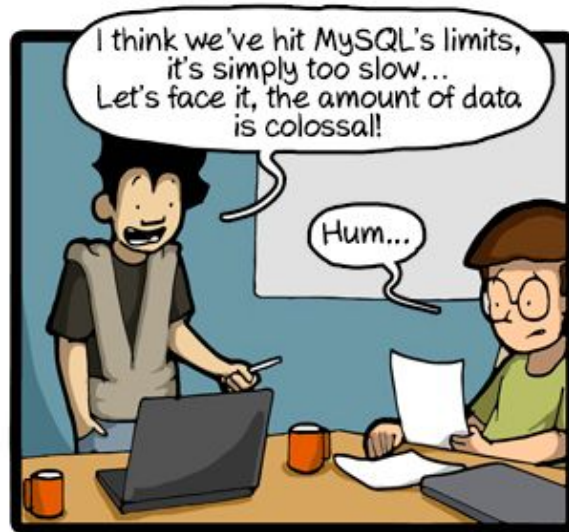
BUT YOU'RE SO



OLD????????

History...

	SQL	MySQL
Creation	1986 (1974)	1995
Last Version	2011 http://modern-sql.com	2015 (5.7)







Application

Framework

ORM

DBAL

SQL

Law of Leaky Abstraction

“All non-trivial abstractions, to some degree, are leaky.”

<http://www.joelonsoftware.com/articles/LeakyAbstractions.html>

J. Spolsky, 11/11/2002

Application

Framework

ORM

DBAL

SQL



YOU'RE DOING IT WRONG.

No matter how hard you try, it is impossible to fax a cat.



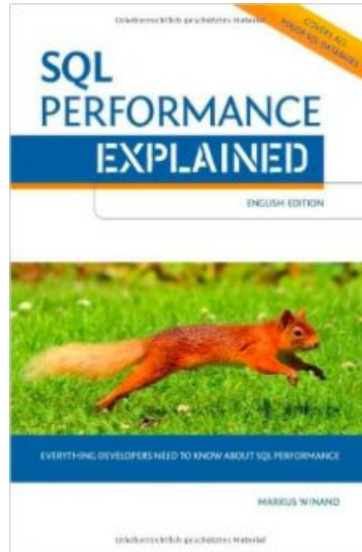
UNDERSTAND SQL

memegenerator.net



USE THE INDEX, LUKE

A guide to database performance for developers



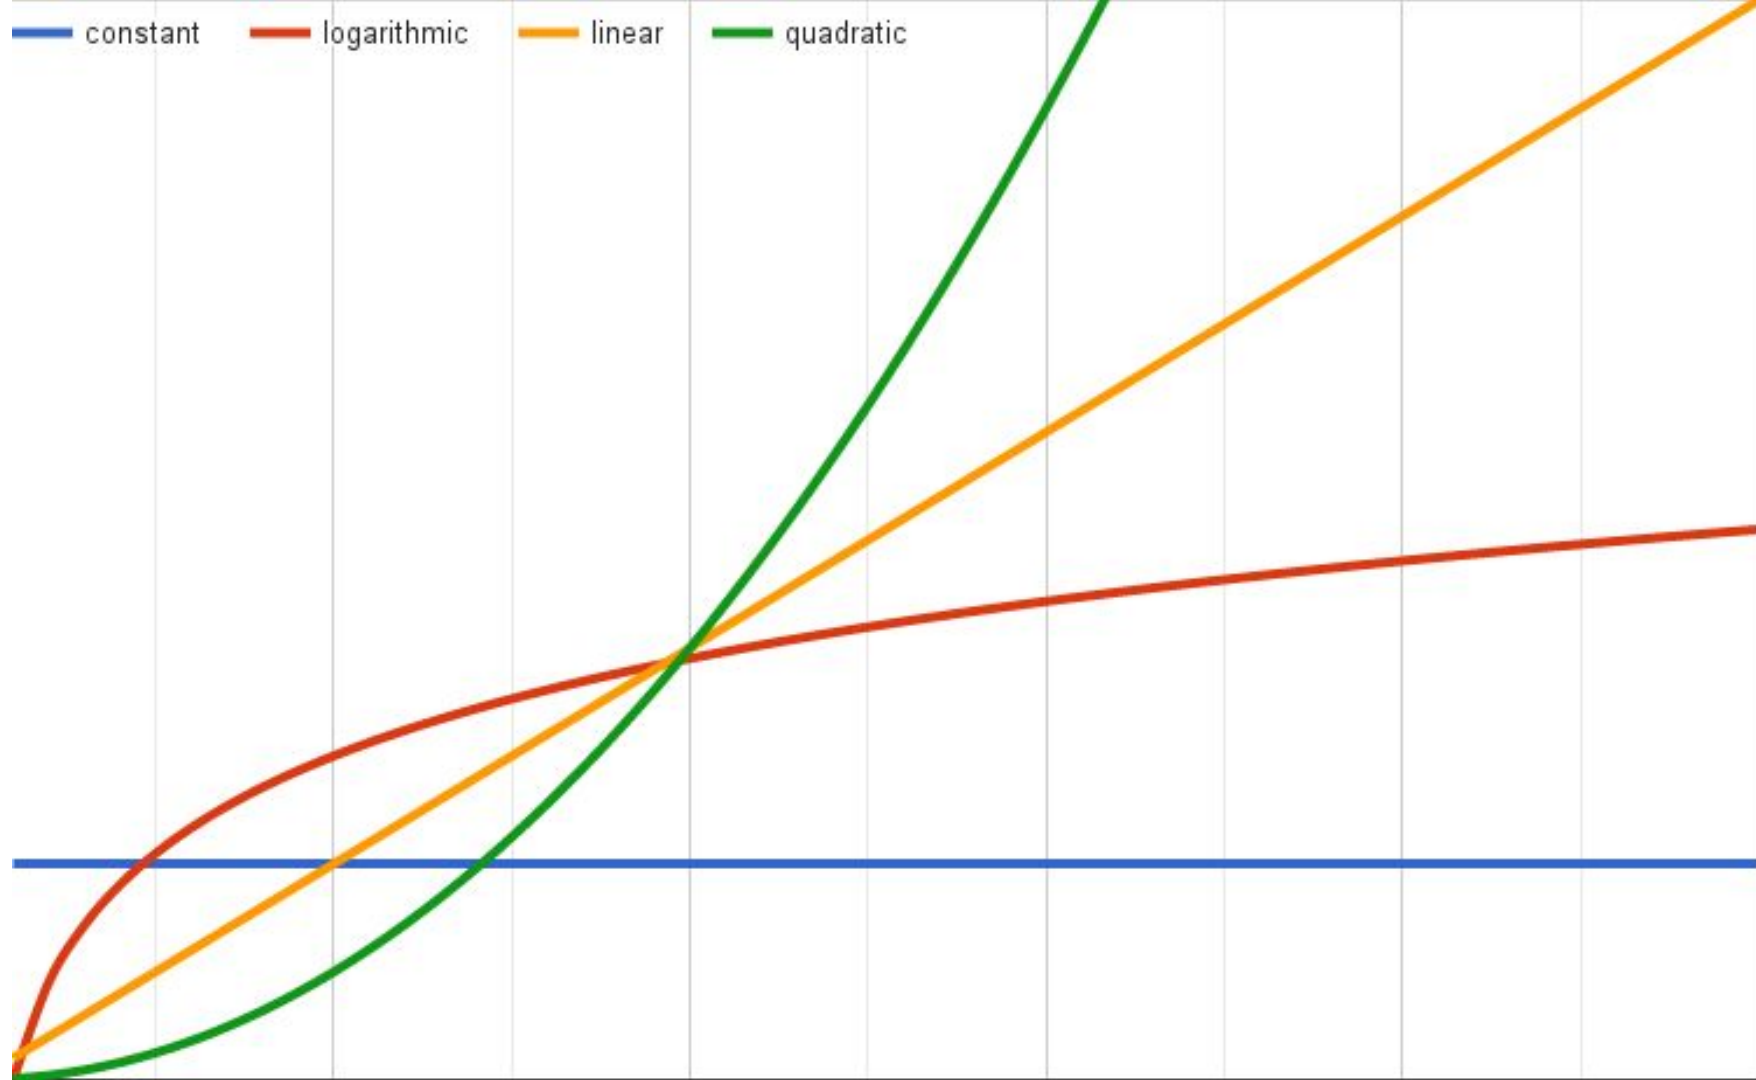
<http://use-the-index-luke.com>

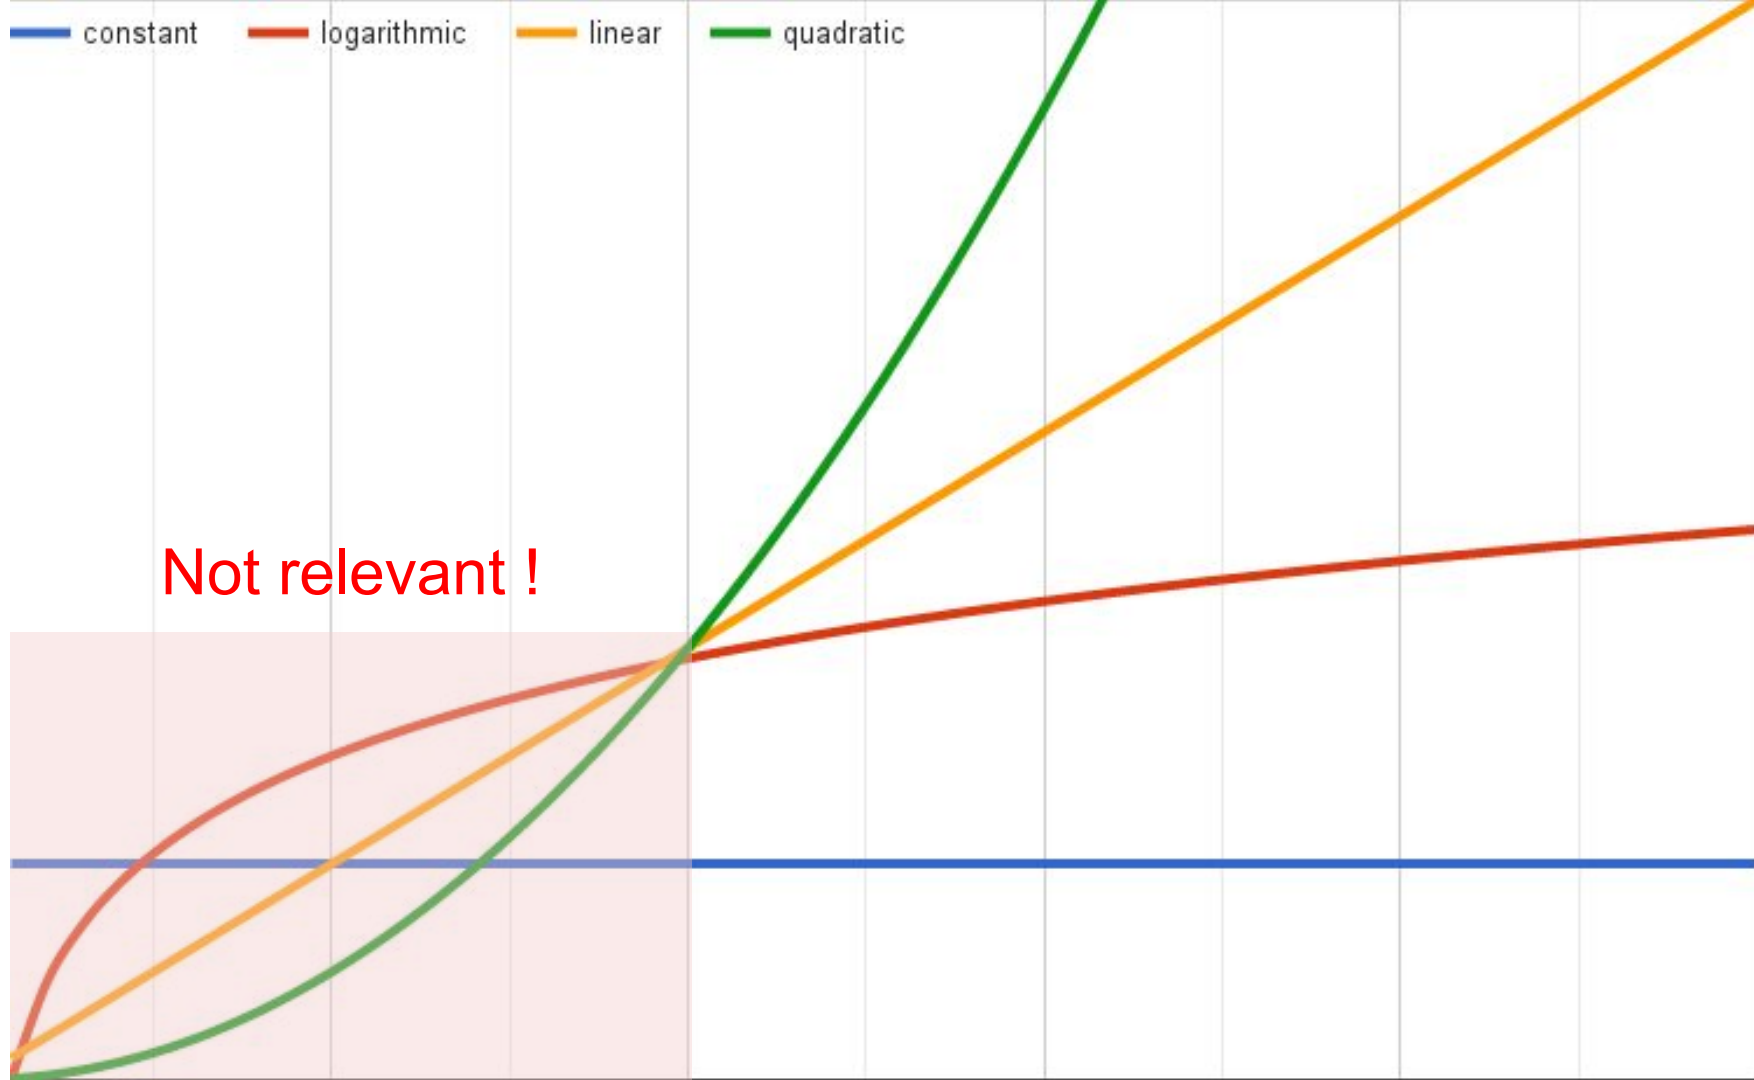


<http://coding-geek.com/how-databases-work/>



WHAT ABOUT COMPLEXITY?





Tables

Back to basics...

Array

	column 0	column 1	column 2	column 3
row 0	Robert	55	manager	USA
row 1	Alex	23	developer	GER
row 2	Jennifer	35	manager	FRA
row 3	Robert	45	CEO	USA
row 4	Charles	32	DBA	UK
...				
row n	Alice	34	developer	ITA

Array

	column 0	column 1	column 2	column 3
row 0	Robert	55	manager	USA
row 1	Alex	23	developer	GER
row 2	Jennifer	35	manager	FRA
row 3	Robert	45	CEO	USA
row 4	Charles	32	DBA	UK
...				
row n	Alice	34	developer	ITA

Full
Table
Scan

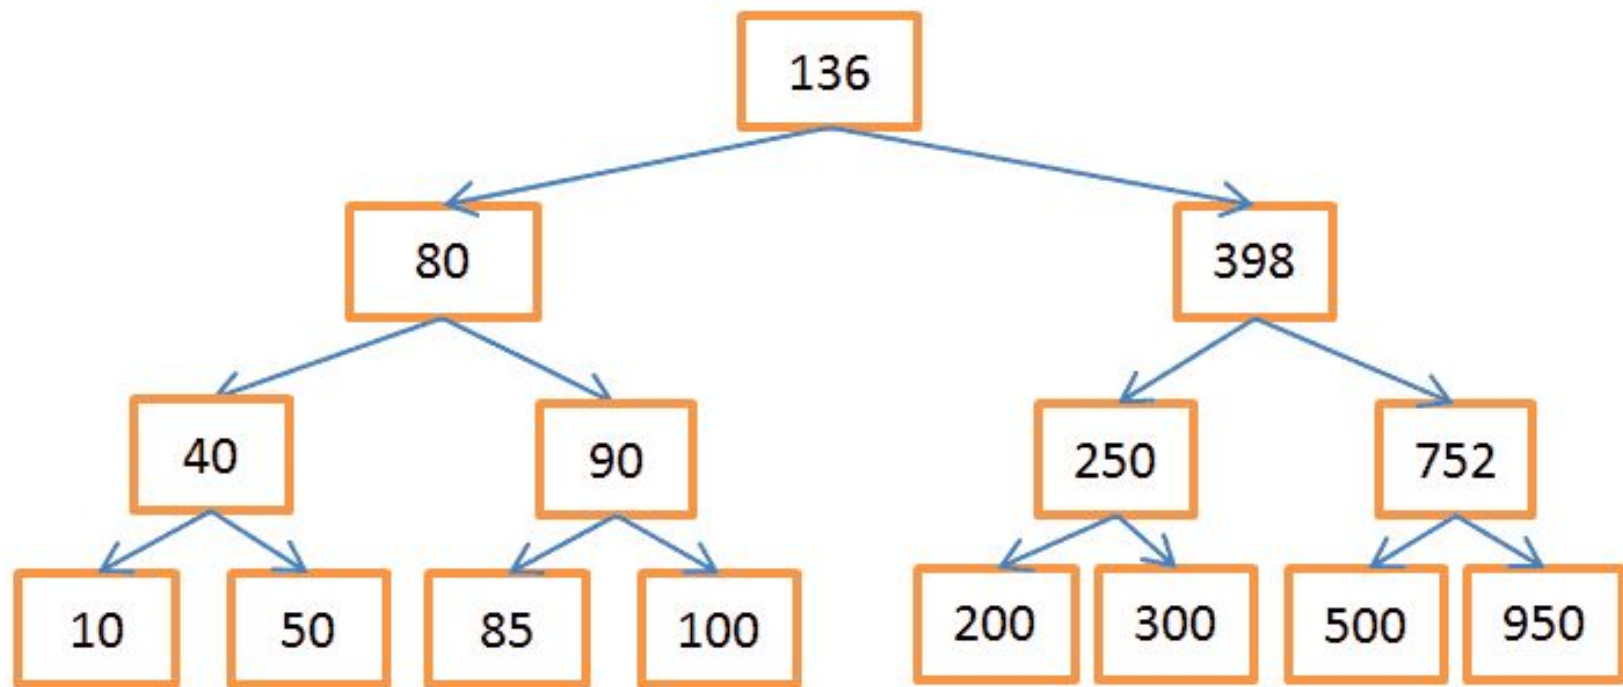
$O(n)$



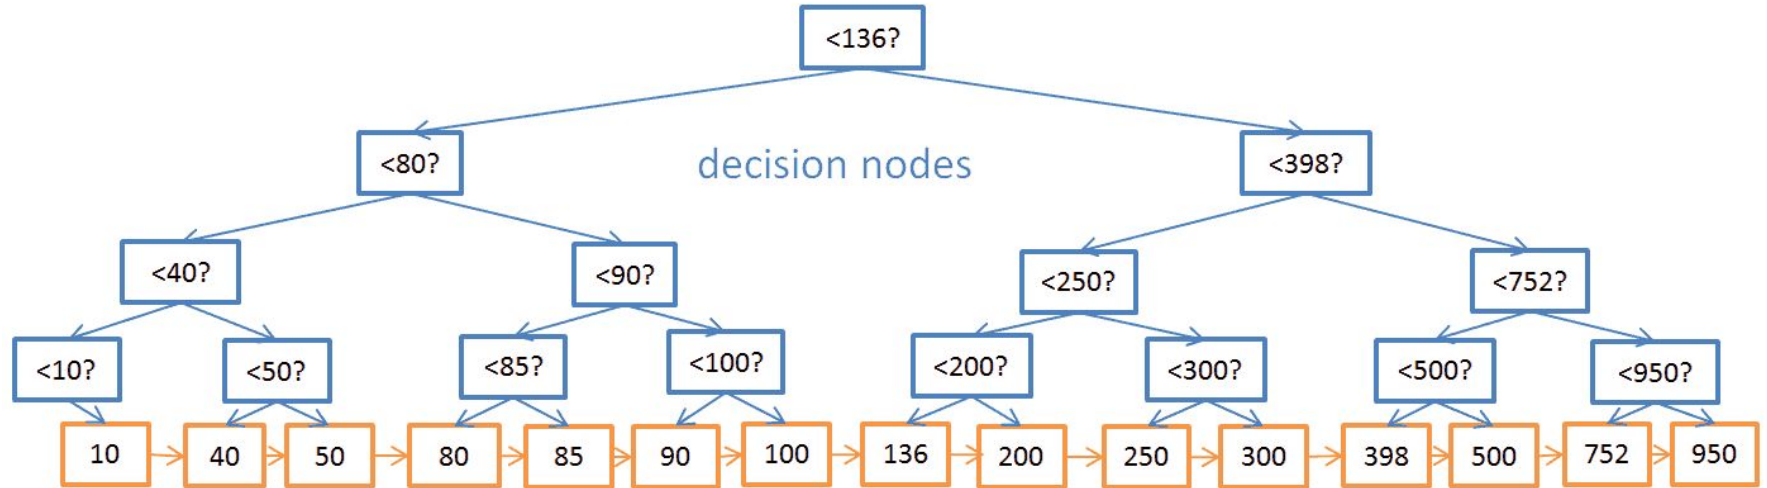
B-Tree

...or not to be...

Binary Search Tree

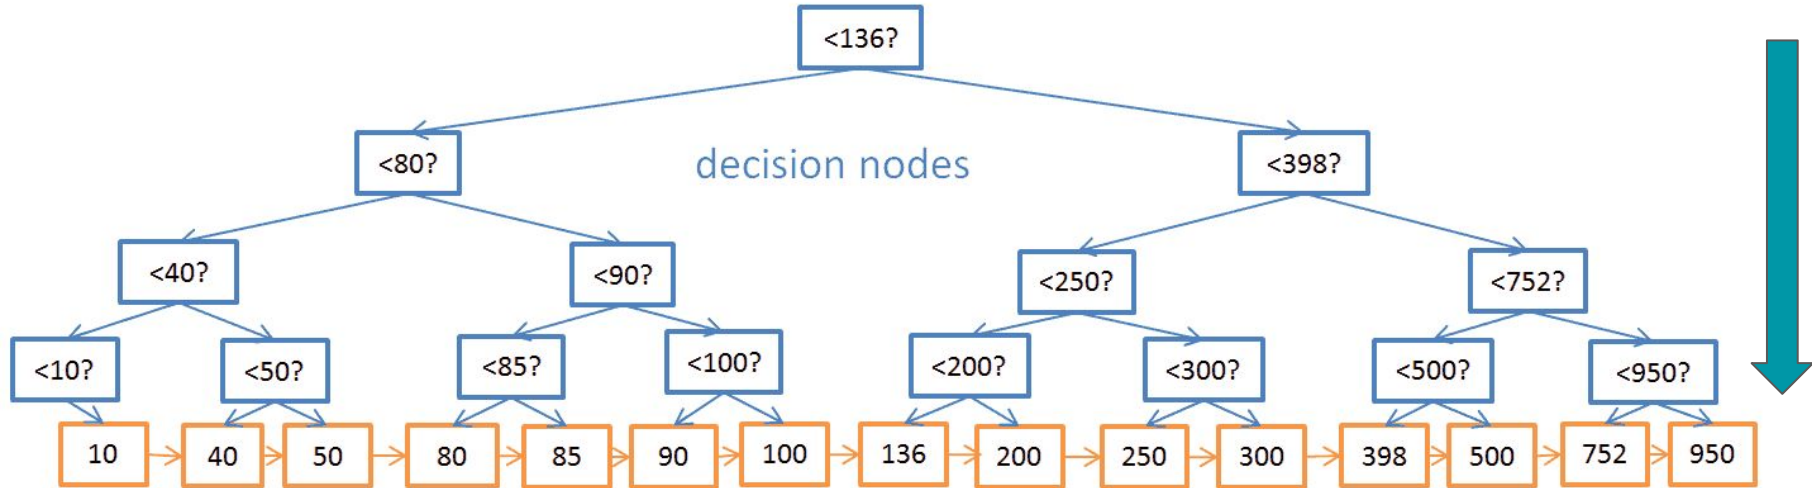


B+ Tree / Database index



B+ Tree / Database index

Index unique scan
 $O(\log n)$



Index Unique Scan $O(\log n)$

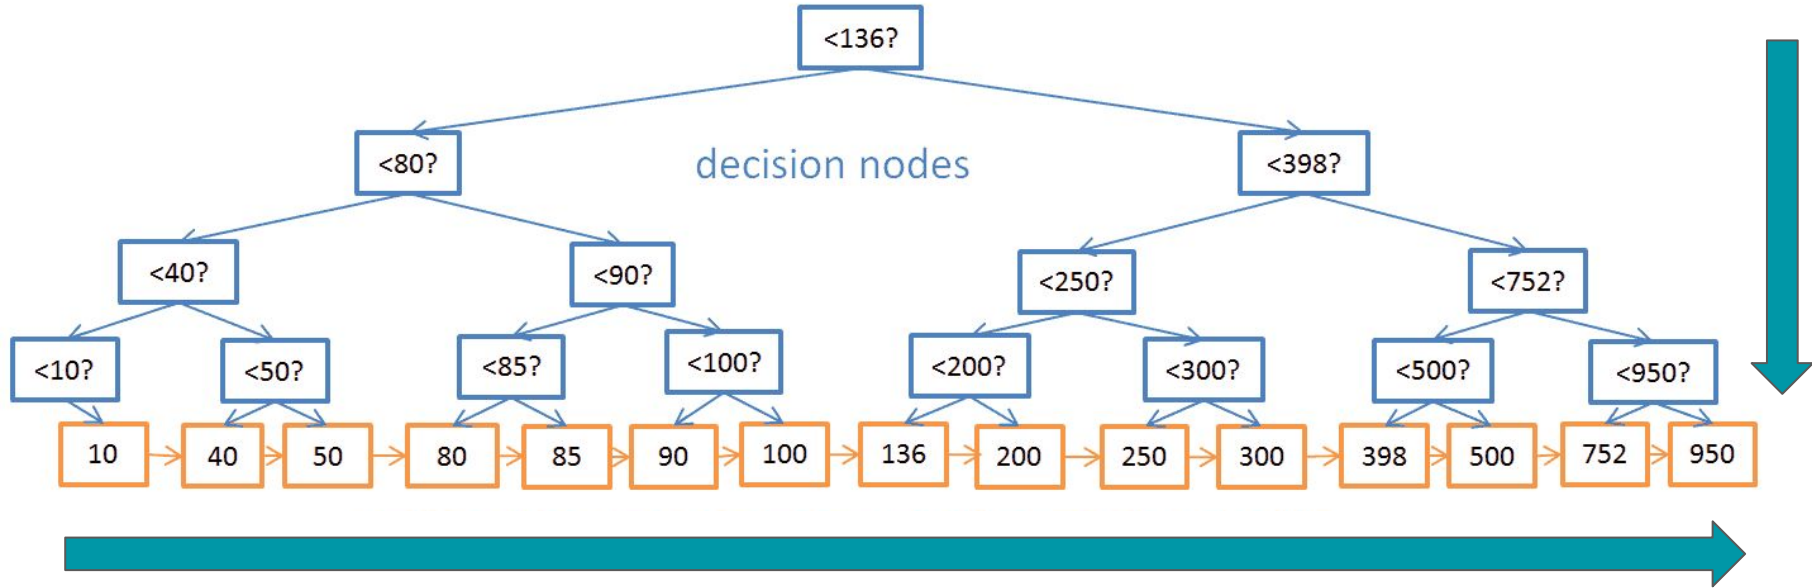
```
SELECT * FROM table WHERE id = ?
```

```
SELECT * FROM table WHERE id != ?
```

```
SELECT * FROM table WHERE id + 10 = ?
```

B+ Tree / Database index

Index unique scan
 $O(\log n)$



Range scan $O(n)$

Range Scan $O(n)$

```
SELECT * FROM table WHERE id BETWEEN ? AND ?
```

```
SELECT * FROM table WHERE id > ? OR data = ?
```

```
SELECT * FROM table WHERE id > 0
```

Text Search

LIKE operator

LIKE 'WI%ND'

WIAW

WIBLQQNPUA

WIBYHSNZ

WIFMDWUQMB

WIGLZX

WIH

WIHTFVZNLC

WIJYAXPP

WINAND

WINBKYDSKW

WIPOJ

WISRGPK

WITJIVQJ

WIW

WIWGPJMQGG

WIWKHLBJ

WIYETHN

WIYJ

LIKE 'WIN%D'

WIAW

WIBLQQNPUA

WIBYHSNZ

WIFMDWUQMB

WIGLZX

WIH

WIHTFVZNLC

WIJYAXPP

WINAND

WINBKYDSKW

WIPOJ

WISRGPK

WITJIVQJ

WIW

WIWGPJMQGG

WIWKHLBJ

WIYETHN

WIYJ

LIKE 'WINA%'

WIAW

WIBLQQNPUA

WIBYHSNZ

WIFMDWUQMB

WIGLZX

WIH

WIHTFVZNLC

WIJYAXPP

WINAND

WINBKYDSKW

WIPOJ

WISRGPK

WITJIVQJ

WIW

WIWGPJMQGG

WIWKHLBJ

WIYETHN

WIYJ

Range Scan $O(n)$

```
SELECT * FROM table WHERE text LIKE 'abc%'
```

```
SELECT * FROM table WHERE text LIKE '%abc%'
```

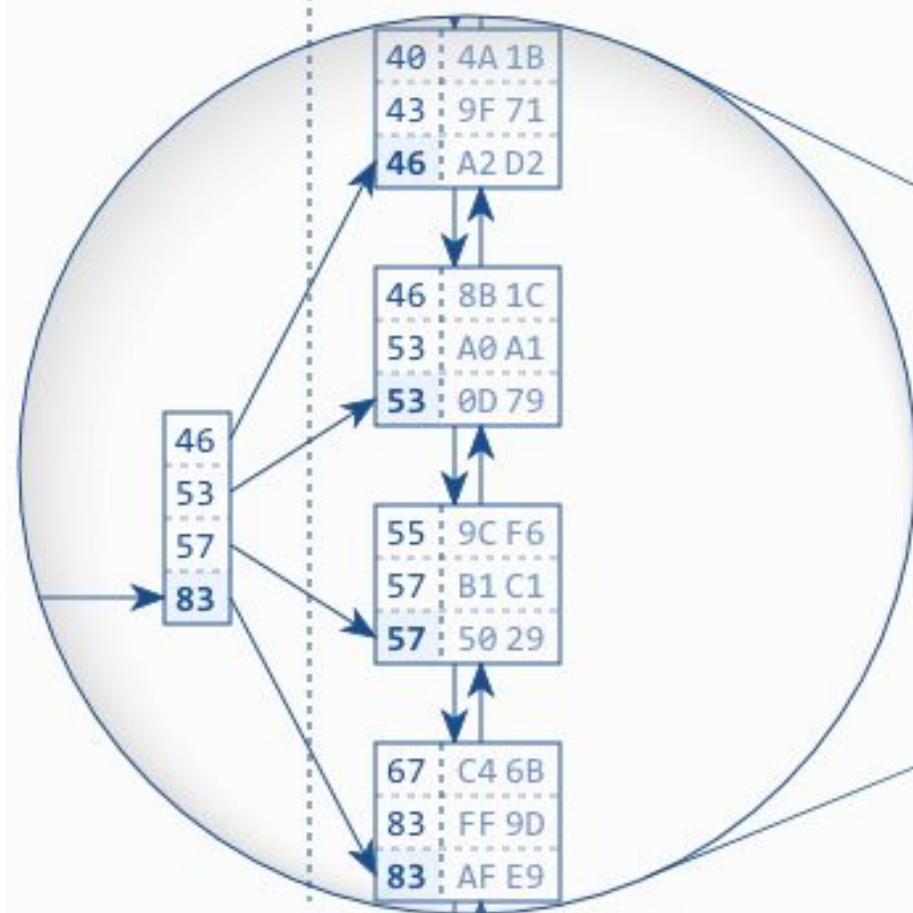
```
SELECT * FROM table WHERE UPPER(text) = 'abc'
```

More nodes!

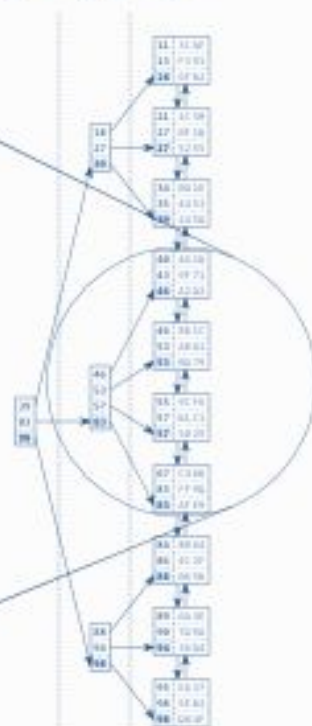
Logarithmic base

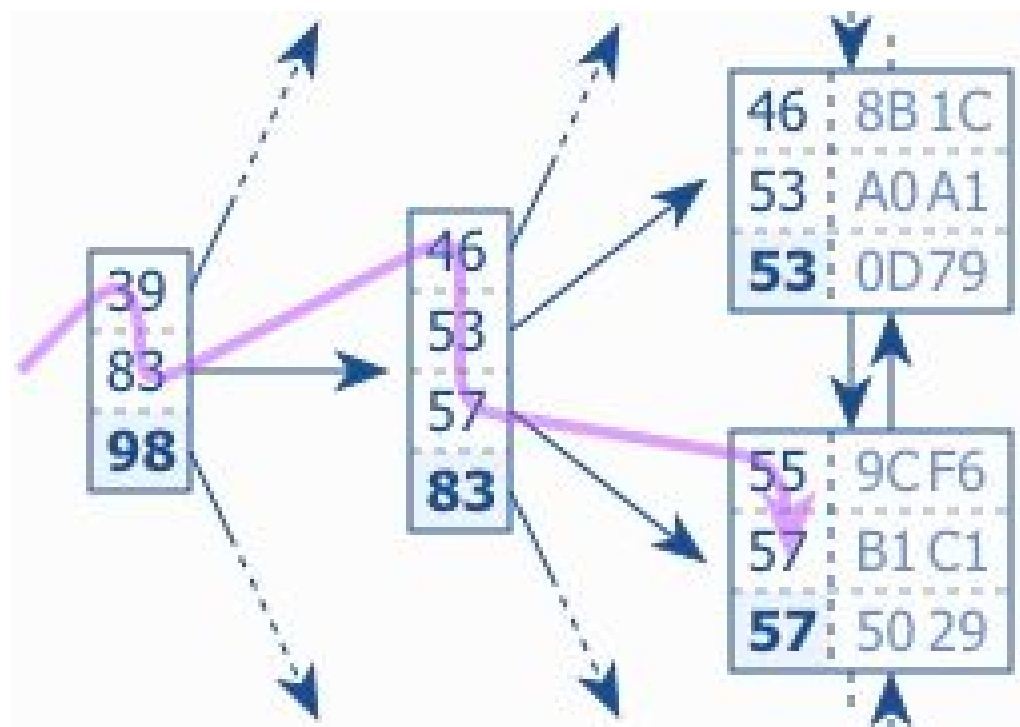
Branch Node

Leaf Nodes



Root Node
Branch Nodes
Leaf Nodes





Data access

The *hidden* cost

Index Leaf Nodes
(sorted)

column 2
ROWID

11	3C AF
13	F3 91
18	6F B2

21	2C 50
27	0F 1B
27	52 55

34	0D 1E
35	44 53
39	24 5D

Table
(not sorted)

column 1
column 2
column 3
column 4

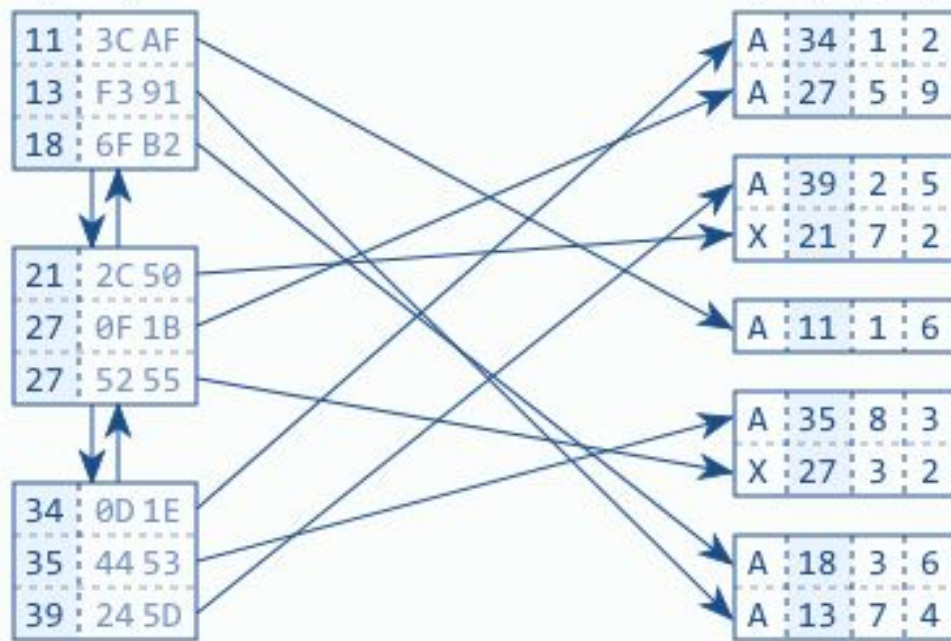
A	34	1	2
A	27	5	9

A	39	2	5
X	21	7	2

A	11	1	6
---	----	---	---

A	35	8	3
X	27	3	2

A	18	3	6
A	13	7	4



Index Leaf Nodes
(sorted)

column 2
ROWID

11	3C AF
13	F3 91
18	6F B2

21	2C 50
27	0F 1B
27	52 55

34	0D 1E
35	44 53
39	24 5D

Table
(not sorted)

column 1
column 2
column 3
column 4

A	34	1	2
A	27	5	9

A	39	2	5
X	21	7	2

A	11	1	6
---	----	---	---

A	35	8	3
X	27	3	2

A	18	3	6
A	13	7	4



Table access by RowID $O(1)$

Clustering

Index-Only scan

Index Only Scan $O(\log n)$

```
SELECT id FROM table WHERE id = ?
```

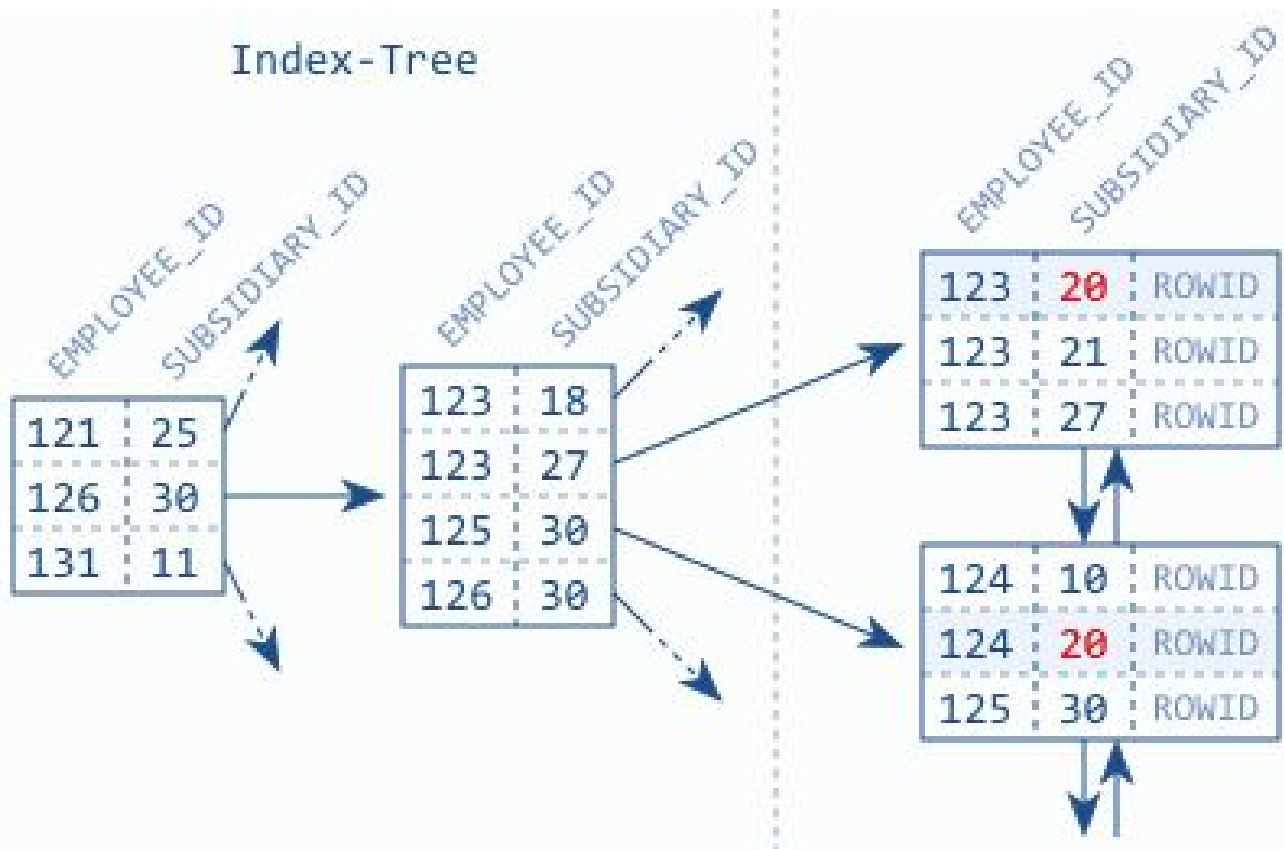
```
SELECT * FROM table WHERE id = ?
```

Multi-column index

Order matters!



Index-Tree



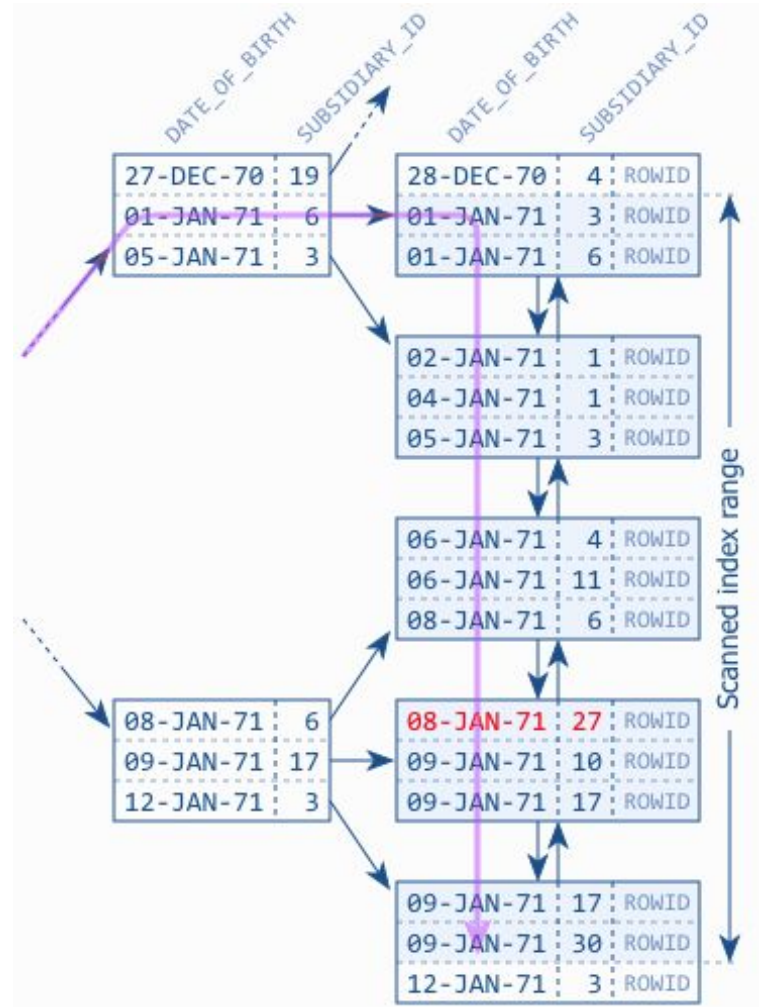
EMPLOYEE_ID, SUBSIDIARY_ID

Range first

WHERE

subsidiary_id = ?

AND date_of_birth BETWEEN ? AND ?

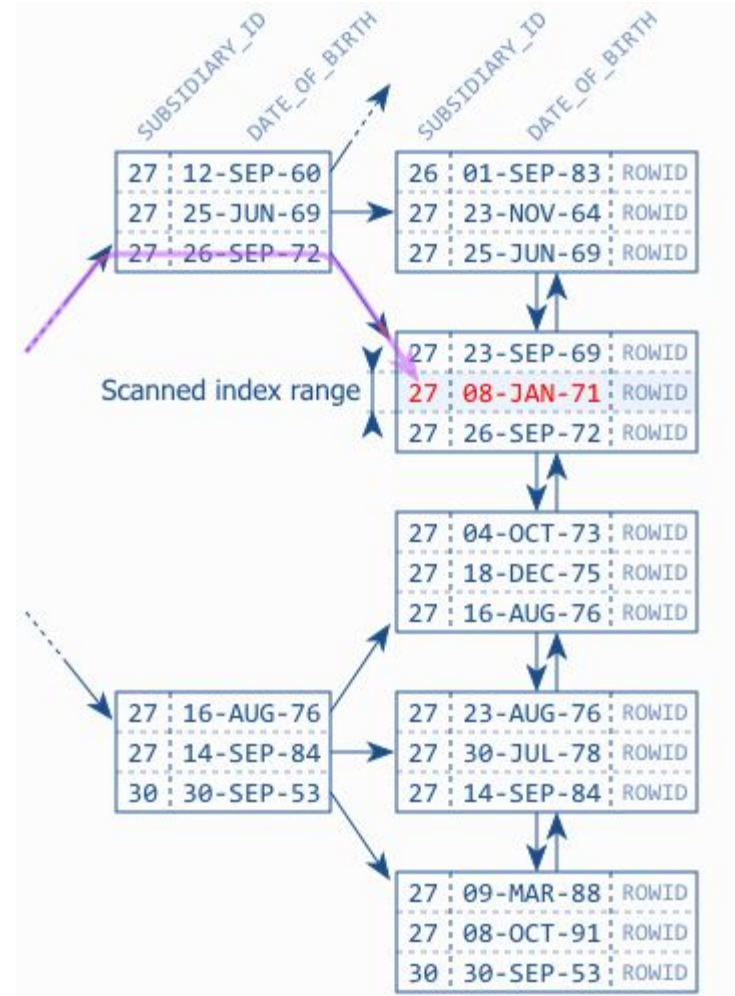


Equality first

WHERE

subsidiary_id = ?

AND date_of_birth BETWEEN ? AND ?



Range Scan $O(n)$

```
SELECT * FROM table WHERE id1 = ?
```

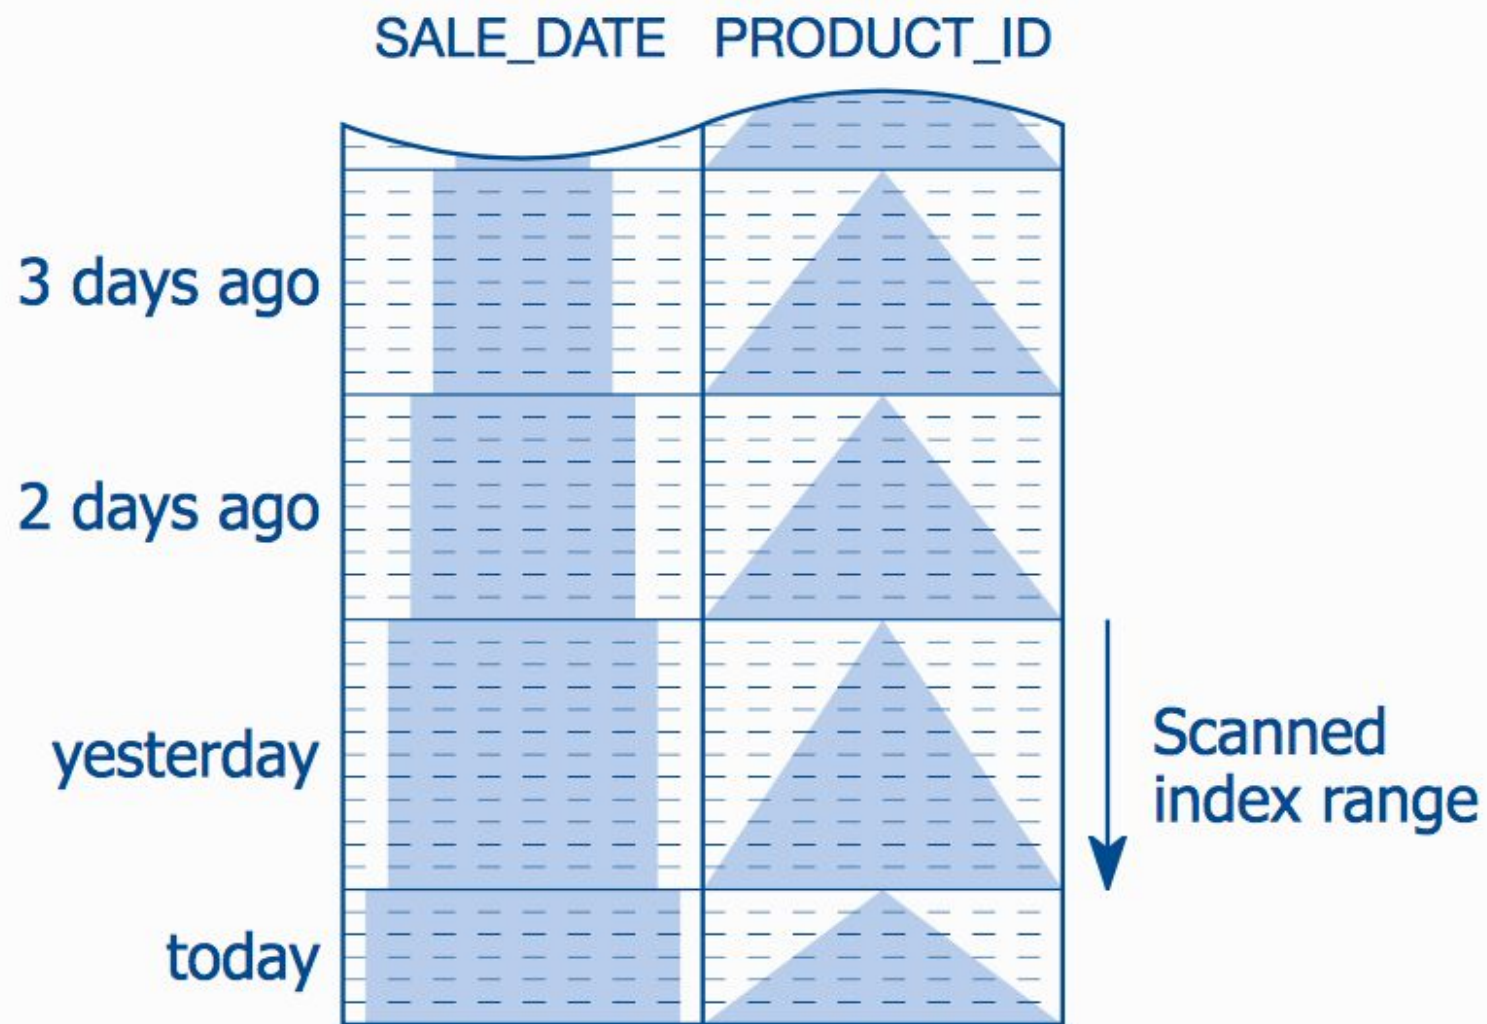
```
SELECT * FROM table WHERE id1 = ? AND id2 > ?
```

```
SELECT * FROM table WHERE id2 = ?
```

```
SELECT * FROM table WHERE id1 > ? AND id2 = ?
```

Sorting

Use the index, luke!



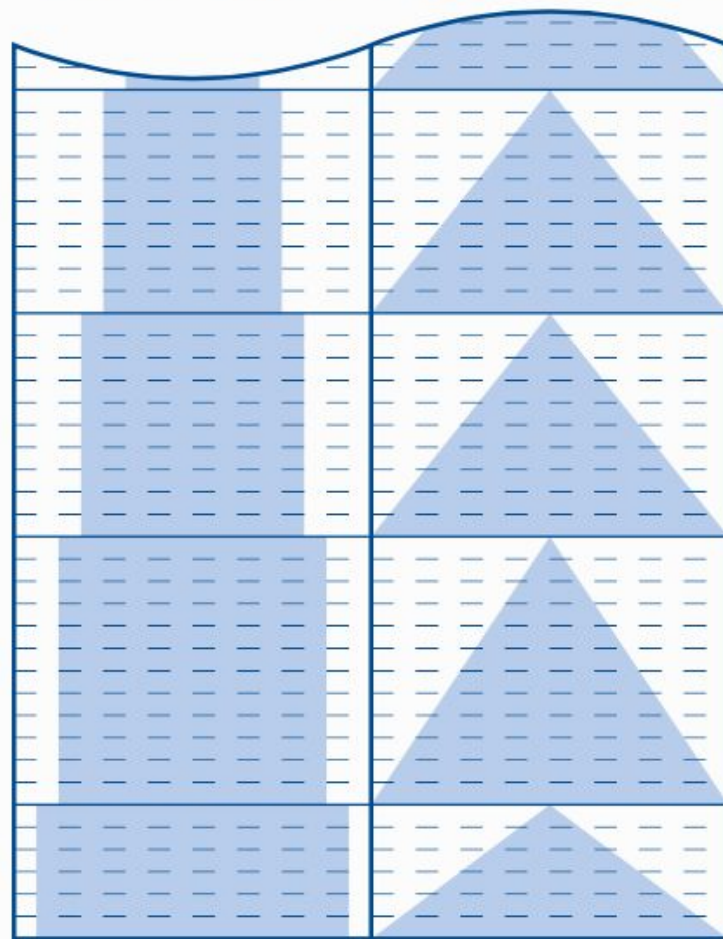
SALE_DATE PRODUCT_ID

3 days ago

2 days ago

yesterday

today



Impossible
index jump

The diagram shows two blue arrows pointing upwards. A red dashed line connects the top of the lower arrow to the bottom of the upper arrow, forming a loop. This indicates a jump in the index that is not possible in a sequential scan.

Grouping

Use the sorting, luke!

Top-N Queries

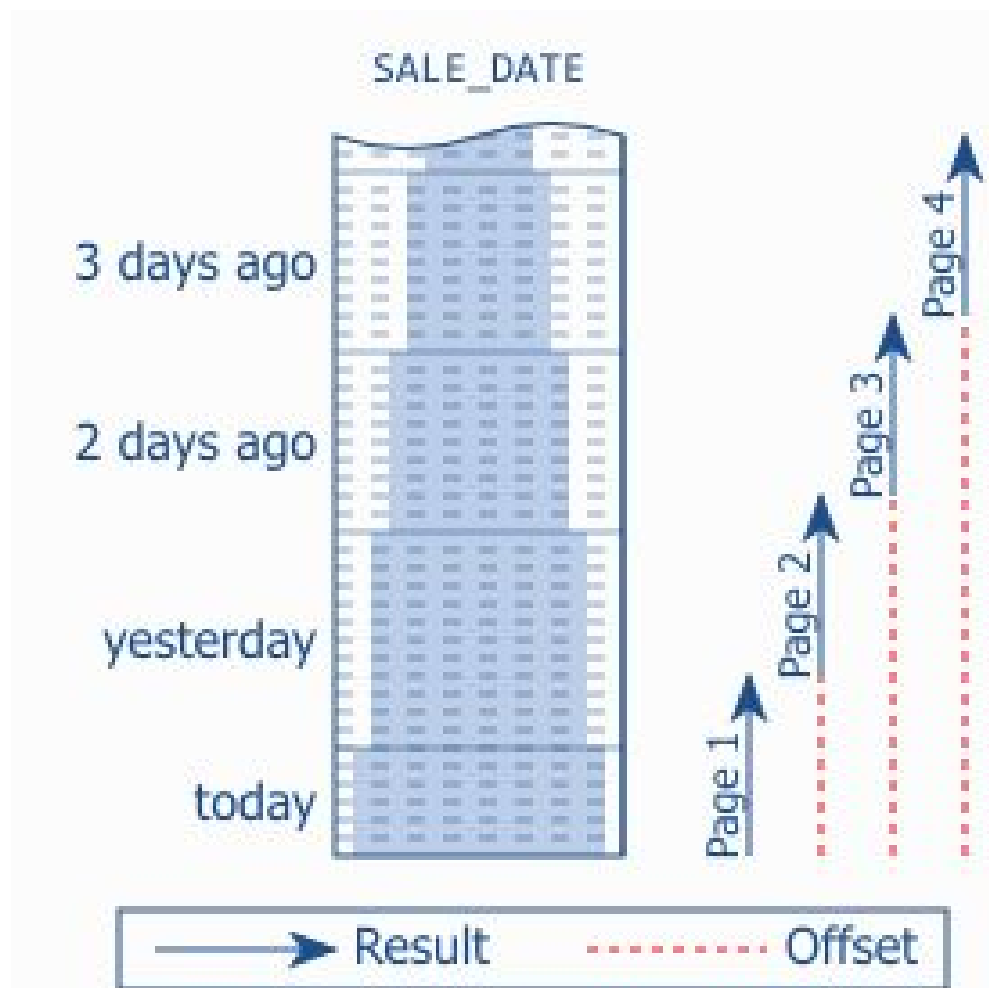
Avoid full table scan!

Offset

Pagination helper



IT'S A TRAP !



BRACE YOURSELF



JOINS ARE COMING

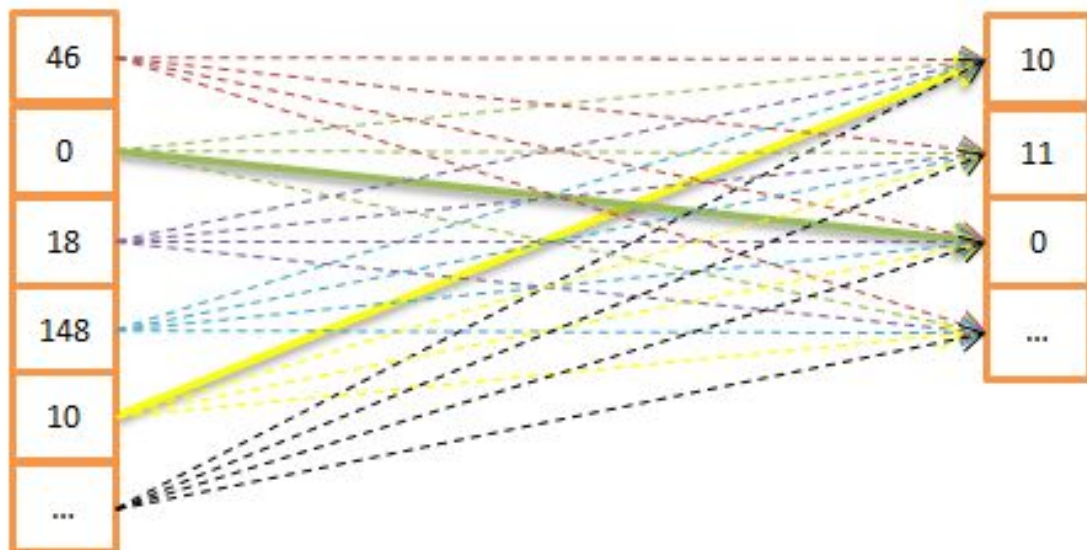
memegenerator.net

Join

Combinatorial power!



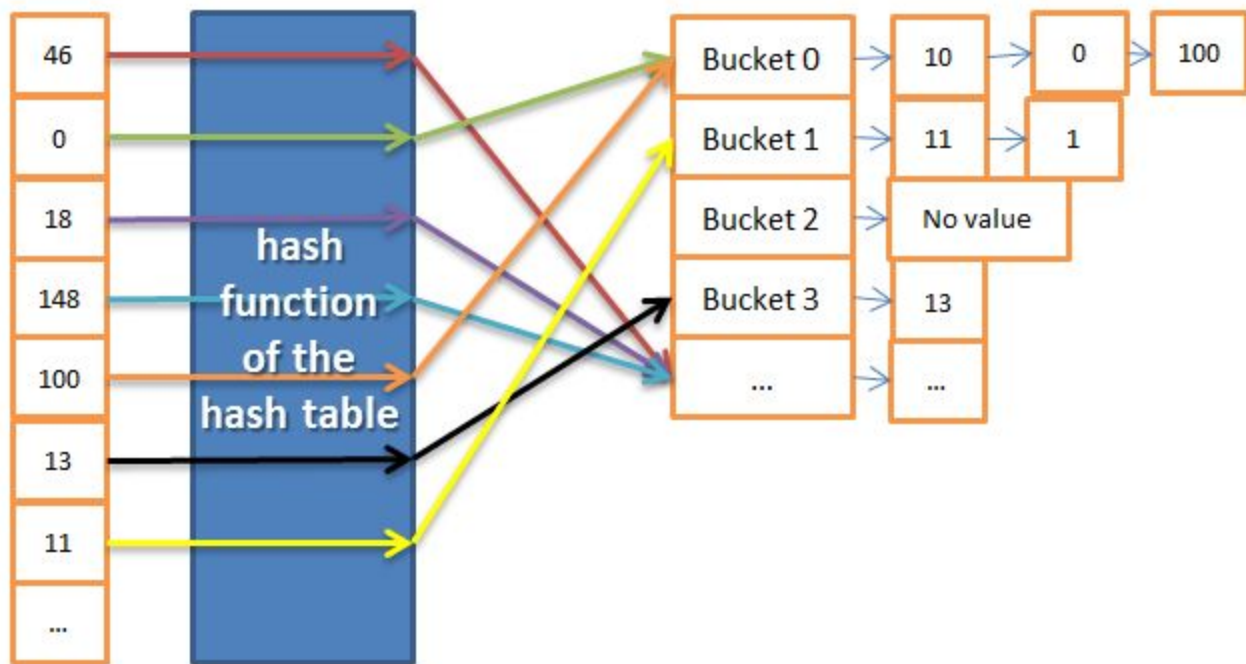
Nested Loop Join



Outer relation

Inner relation

Hash Join



Outer relation

Inner relation (in-memory hash table)

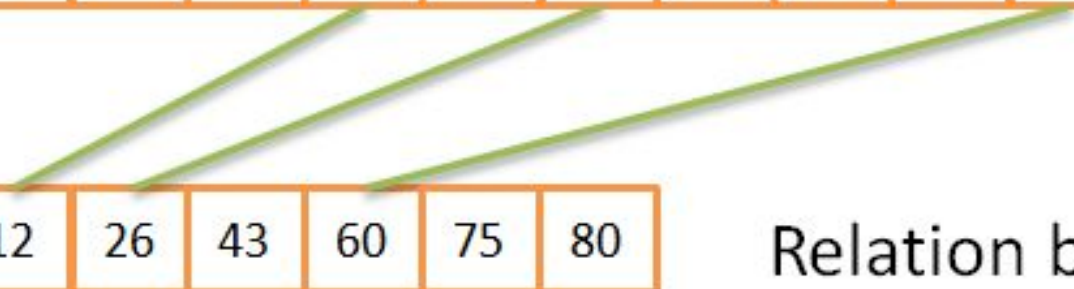
Merge Join

1	6	10	11	12	24	26	30	31	40	60	70
---	---	----	----	----	----	----	----	----	----	----	----

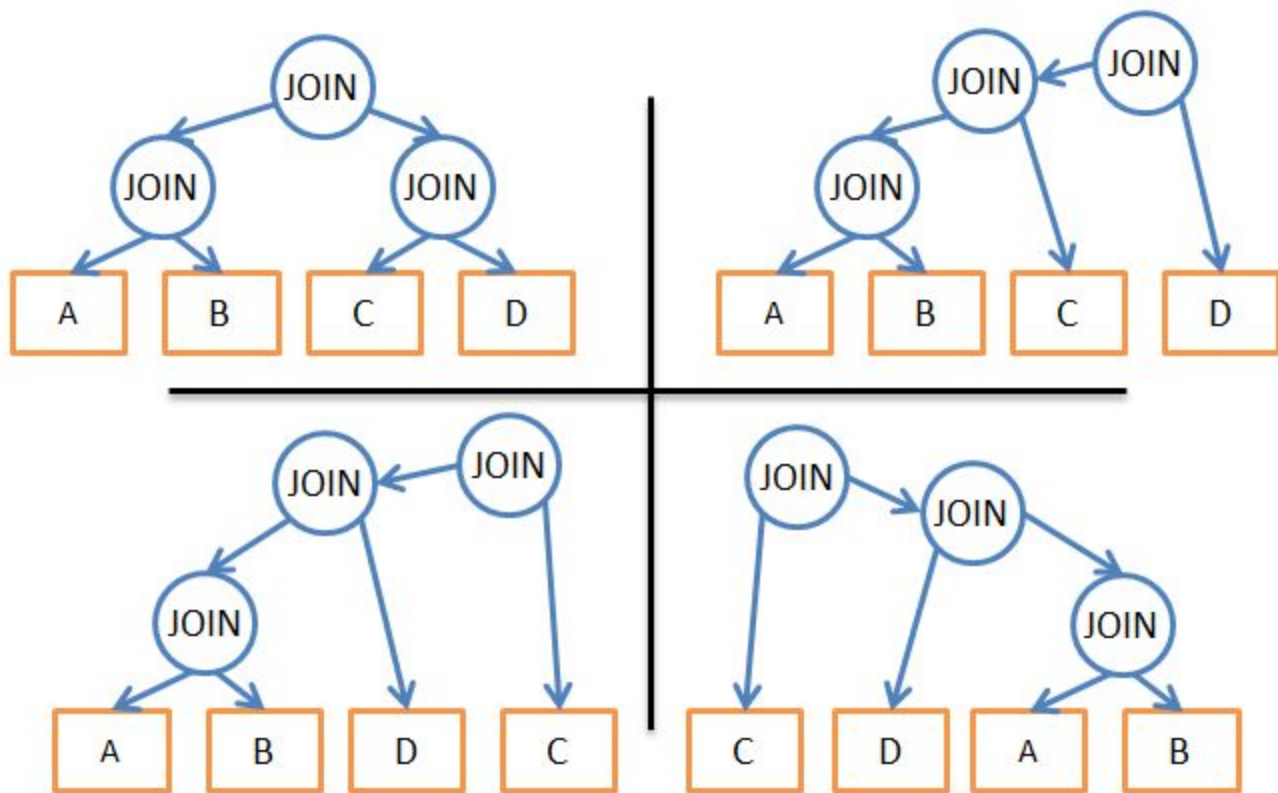
Relation a

3	12	26	43	60	75	80
---	----	----	----	----	----	----

Relation b



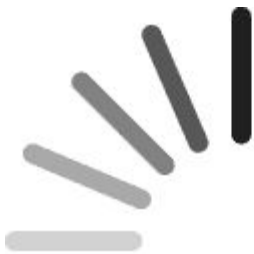
Join ordering problem



Statistics

Best execution plan





...

1 Query

≈

1 Index

Speed vs (Memory and Speed update)

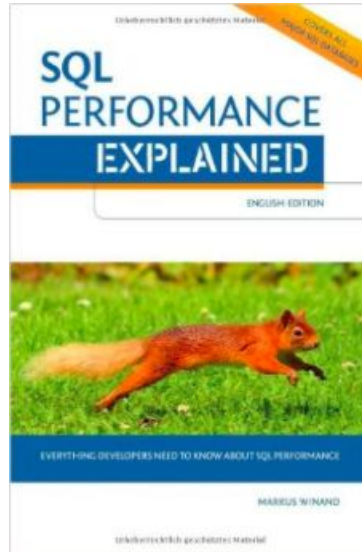
And now?

- Slow Log
- EXPLAIN



USE THE INDEX, LUKE

A guide to database performance for developers



<http://use-the-index-luke.com>



<http://coding-geek.com/how-databases-work/>



Thanks!